

**Question 1 Skipped**

^

Your organization uses Google Kubernetes Engine (GKE) for its microservices-based application. The number of services is expected to grow significantly over the next two years. What would be the best approach to ensure manageability and operational efficiency as the number of services increases?

Increase the size of the GKE cluster nodes to accommodate the growing number of services.

Migrate to Compute Engine instances to reduce the complexity of managing Kubernetes.

Implement Cloud Logging for better log management.

Correct answer



Implement a Service Mesh using Istio for fine-grained control and observability.

Question 2 Skipped ^

A multinational retail company wants to implement a new e-commerce platform that can handle high traffic and provide real-time recommendations to users. The platform should also be able to track user behavior and provide insights into customer preferences. The company wants to use the Google Cloud Platform for this implementation. Which of the following options would be the best solution for this requirement?

Correct answer



- Use Cloud SQL to store customer data, Cloud Pub/Sub to ingest customer data, and Cloud Dataflow to process customer analytics.

- Use Cloud Firestore to store customer data, Cloud Functions to process customer data, and Cloud Datalab to store customer analytics.

- Use Cloud Bigtable to store customer data, Cloud Pub/Sub to ingest customer data, and Cloud Functions to process customer analytics.

- Use Cloud Storage to store customer data, Cloud Functions to process customer data, and Cloud Datastore to store analytics data.

Question 3 Skipped ^

You're creating a custom VPC in Google Cloud with several subnets. One of your subnets requires communication with an on-premises network via Cloud VPN. What should you consider when designing this scenario?

Correct answer



- Ensure the CIDR range of the VPC subnet does not overlap with the on-premises network.

- Assign public IP addresses to all instances in the VPC subnet.

- Create a Shared VPC instead of a custom VPC.

- The CIDR range of the on-premises network and the VPC subnet should be the same for easy routing.

**Question 4 Skipped**

^

As a cloud architect, you are designing a secure architecture for a client who needs to limit the use of external IP addresses on their Compute Engine instances. The client wants to ensure that only specific, approved instances can be assigned external IP addresses. What method would be the most appropriate for enforcing this requirement?

- Use VPC Service Controls to restrict instances from obtaining external IP addresses.

Correct answer



- Use an organization policy to restrict external IP addresses and apply it to the project.

- Use firewall rules to block outbound traffic from instances without approved external IPs.

- Use Shared VPC to restrict instances from obtaining external IP addresses.

 **Question 5 Skipped**

^

Your company has a robust e-commerce application running on Compute Engine instances and using Cloud Storage. As the business expands globally, you notice latency issues affecting customers in different regions. What approach should you consider to improve your solution?

- Migrate from Cloud Storage to Persistent Disk for faster I/O operations.
- Increase the size of the Compute Engine instances to enhance the application's performance.
- Move the application to App Engine to leverage its automatic scaling features.
- Implement a Content Delivery Network (CDN) using Cloud CDN to cache static content closer to users.  


Correct answer

 **Question 6 Skipped**

^

When planning your migration, you find out that some members of the network management team will need to be able to manage all network components, but other team members will only need read access. What mechanism should you use to control this?

**Virtual Private Networks**

**Firewall rules**

**Virtual Private Clouds**

**Correct answer**



**IAM roles**

● Question 7 Skipped ^

As a cloud architect, you are working with a client who operates a popular web service. Their traffic patterns are highly variable, and they need a deployment strategy that will scale in response to the load on their system. The client wants to deploy a Managed Instance Group (MIG) with a minimum of 5 instances and a maximum of 50 instances. Which of the following strategies should you use to meet these requirements?

- Deploy the MIG with a minimum of 5 instances, but do not specify a maximum number of instances.
- Deploy the MIG with Autoscaler, and use Instance Templates to specify the maximum number of instances.
- Deploy the MIG and configure it with an Autoscaler, setting the target utilization level to match the desired load.
- Deploy the MIG without Autoscaler, manually adding instances as needed.

Correct answer



 **Question 8** Skipped ^

You are creating a VPC network to host a set of compute resources for a multi-tier web application in Google Cloud. How should you configure the network to optimize security and manageability?

- Use automatic mode to create the VPC network.
- Assign public IP addresses to all compute instances to facilitate direct access.
- Use Shared VPC to host all the compute resources.

Correct answer



- Create separate subnets for each tier of the application.

 **Question 9** Skipped 

Your company's data science team anticipates a significant increase in the number and complexity of machine learning (ML) models they plan to train over the next year. Currently, the team uses Compute Engine instances for training. What would be the best approach to accommodate this increase?

- Increase the disk size for the Compute Engine instances to store more training data.

**Correct answer**



- Use AI Platform Training to offload and scale model training tasks.

- Upgrade the Compute Engine instances to have more CPUs and more memory.

- Use Cloud Functions to train individual models in response to events.

Question 10 Skipped

^

As a cloud architect, you are working with an e-commerce client who expects a significant increase in traffic to their application hosted on Google Kubernetes Engine (GKE) during an upcoming sales event. To manage the increased load, the client wishes to enable autoscaling for their application. However, they want to ensure that any newly created pods during autoscaling are fully ready to handle traffic before being included in the service. Which of the following solutions should you implement?

Correct answer



- Configure a readiness probe in your pod specification.
- Configure a startup script in the pod specification to delay service registration.
- Configure a liveness probe in your pod specification.
- Enable Cluster Autoscaler on your GKE cluster.

### Question 11 Skipped

^

As a cloud architect, you are working with an organization to structure their BigQuery permissions. The company has three teams: data analysts, data scientists, and data engineers. The data analysts should be able to run SQL queries, the data scientists need to create and run machine learning models in BigQuery, and the data engineers should have full control over BigQuery resources. What is the most appropriate IAM role assignment for these teams?

**Assign the roles/bigquery.dataEditor role to the data analysts, roles/bigquery.mlUser role to the data scientists, and roles/bigquery.admin role to the data engineers.**

- Assign the roles/bigquery.dataEditor role to the data analysts, roles/bigquery.mlUser role to the data scientists, and roles/bigquery.admin role to the data engineers.

- Assign the roles/bigquery.jobUser role to the data analysts, roles/bigquery.mlUser role to the data scientists, and roles/bigquery.dataEditor role to the data engineers.

Correct answer



- Assign the roles/bigquery.user role to the data analysts, roles/bigquery.mlUser role to the data scientists, and roles/bigquery.admin role to the data engineers.

**Assign the roles/bigquery.dataViewer role to the data analysts, roles/bigquery.mlUser role to the data scientists, and roles/bigquery.dataOwner role to the data engineers.**

- Assign the roles/bigquery.dataViewer role to the data analysts, roles/bigquery.mlUser role to the data scientists, and roles/bigquery.dataOwner role to the data engineers.

● **Question 12 Skipped**



As a cloud architect, you are planning to run stateful applications in Kubernetes Engine. What should you use to support stateful applications?

**DeamonSet**

**StatefulPods**

**Correct answer**



**StatefulSets**

**ReplicaSet**

**Pods**



### Question 13 Skipped



As a cloud architect, you are working with a rapidly expanding startup that uses Google Cloud. They have been manually creating individual Google user accounts for each employee, but now they are scaling and need a more efficient and secure way to manage access to Google Cloud resources.

Which of the following options should you recommend?

Correct answer



- Set up a Cloud Identity or G Suite domain, then use Google Cloud Directory Sync to synchronize the users to Google Cloud.

- Continue creating individual user accounts for each employee manually.

- Create service accounts for each employee and distribute the private keys.

- Use the Google Cloud IAM service to create a custom IAM role for each employee.

Question 14 Skipped ^

As a cloud architect, you are assisting a client to transfer a large set of files from their local data center to Cloud Storage. The client has asked to minimize the costs associated with network egress. They have a Cloud Interconnect connection set up and want to transfer files over this connection. What is the most appropriate gsutil command to use for this purpose?

- `gsutil cp -r ./local_directory gs://bucket_name`
- `gsutil -o "GSUtil:parallel_composite_upload_threshold=150M" cp -r ./local_directory gs://bucket_name`
- `gsutil -o "Boto:parallel_composite_upload_threshold=150M" cp -r ./local_directory gs://bucket_name`
- `gsutil -m cp -r ./local_directory gs://bucket_name`

Correct answer



**Question 15 Skipped**

^

Your organization is planning to create a custom VPC network on Google Cloud for deploying a multi-tier application. The application includes frontend servers, backend servers, and database servers, each of which requires a separate subnet. What considerations should you keep in mind while creating this VPC?

- Use the same CIDR block for all the subnets.

Correct answer



- Assign one large CIDR block to the VPC and divide it into smaller CIDR blocks for each subnet.

- Use Cloud NAT for the frontend servers to connect with the internet.

- Use Shared VPC to host all the tiers of the application.



### Question 16 Skipped



As a cloud architect, you are responsible for preparing migration strategy. Your company needs to mount a shared filesystem to Compute Engine instances. Which GCP service should you use?

Correct answer



- Cloud Filestore

- Cloud Firestore

- Cloud Storage

- Local SSD

● Question 17 Skipped ^

Your company wants to use Google Cloud to host a static website that will receive high amounts of traffic. The website must be highly available and must load quickly for users around the world. How would you host the website in Google Cloud to meet these requirements?

- Host the website on a single Compute Engine virtual machine and use a global load balancer to distribute traffic.

Correct answer



- Host the website on Cloud Storage and use Cloud CDN to distribute traffic.

- Host the website on multiple Compute Engine virtual machines in different regions and use a regional load balancer to distribute traffic.

- Host the website on a single App Engine Standard environment instance.



### Question 18 Skipped



As a cloud architect, you are working with a company to deploy their multi-tier application on Google Kubernetes Engine (GKE). The application includes a front-end service that needs to maintain high availability. The company has requested that at least five replicas of the front-end service run at all times to meet their availability requirements. Which field in the Kubernetes deployment manifest should you use to specify this requirement?

`spec.pods`

`metadata.replicas`

`metadata.scale`

**Correct answer**



`spec.replicas`

● Question 19 Skipped ^

Your company is developing a new mobile application that needs to store user data in the cloud. The data must be highly available and durable, and the application must be able to read and write data even when the device is offline. How would you store the user data in Google Cloud to meet these requirements?

Correct answer



- Use Cloud Firestore in Native mode to store the user data.
- Use Cloud Bigtable to store the user data.
- Use Cloud Firestore in Datastore mode to store the user data.
- Use Cloud Datastore to store the user data.

## Question 20 Skipped



Your customer is planning to run a highly available and scalable web application in Google Cloud.

The customer has the following requirements:

- the application must be easily deployable and manageable
- the application must be highly available and recover from failures automatically
- the application must be able to handle incoming traffic spikes and scale dynamically
- the application must be secure and protect against common web attacks

Which Google Cloud service should you recommend to secure and protect the application against common web attacks?

Google Kubernetes Engine

Cloud Functions

App Engine

**Correct answer**



Cloud Armor

● Question 21 Skipped

^

As a cloud architect, you are working with a global gaming company that needs a database for their new real-time, multi-player game. They require low latency, high transaction throughput, and the ability to scale to millions of users worldwide. Which of the following approaches would be the most effective way to use Google Cloud Spanner to meet these requirements?

Correct answer



- Use Cloud Spanner's multi-region configurations and configure the game to direct traffic based on user location.
- Set up multiple Cloud Spanner instances in each region where the game has users, and manually replicate data between instances.
- Create a single-region Cloud Spanner instance for the lowest possible latency and manually partition the database by user location.
- Create a single Cloud Spanner instance and rely on Google's global network for low latency.



## Question 22 Skipped



Your compliance team is concerned about using a public cloud service because other companies will be running their systems in the same cloud. You assure them that your company's resources will be isolated and inaccessible to others. Which resource is used for this purpose?

Correct answer



- Virtual Private Clouds - VPCs

- Cloud VPN

- CIDR blocks

- Cloud Interconnect

● Question 23 Skipped ^

You have a web application operating on a Managed Instance Group, which receives a high volume of requests every minute. Your objective is to apply patches to the application without reducing the number of instances in the MIG. What action should you take?

- You should deploy the update in a new Managed Instance Group and add it as a backend service to the existing production Load Balancer. Then remove the old Managed Instance Group from the Load Balancer backend and remove the group.

Correct answer



- You should perform a rolling-action start-update with `max-unavailable` set to 0 and `max-surge` set to 1.

- You should perform a rolling-action start-update with `max-unavailable` set to 1 and `max-surge` set to 0.

- You should update the existing Managed Instance Group to point to a new instance template containing the updated version. Terminate all existing instances in this group and wait until they are all replaced by new instances created from the new template.

Question 24 Skipped ^

You run a small startup and want to optimize costs in GCP. As a cloud architect, you need to research resource consumption charges and provide a summary of your expenses. You want to do it in the most efficient way. What should you do?

Correct answer



- You should attach labels to resources to reflect the purpose. Then, export Cloud Billing data into BigQuery, and analyze it with Data Studio.

- You should use Google Cloud Recommender to see if expenses can be reduced.

- You should assign tags to resources to reflect the purpose. Export Cloud Billing data into Cloud SQL, and analyze it with Data Studio.

- You should rename resources to reflect the purpose. Write a Python script to analyze resource consumption.

 **Question 25** Skipped 

As a cloud architect, you are working with a cost-conscious client who regularly runs large, complex queries on BigQuery. They want to understand how they can estimate the cost of running a BigQuery query before they execute it to avoid unexpected expenses. Which of the following is the most suitable solution for this?

- Use Google Cloud Pricing Calculator by entering the estimated amount of data processed by the query.

**Correct answer**



- Preview the query in the BigQuery web UI, which displays an estimate of the amount of data the query will process.

- Use Cloud Monitoring to estimate the cost based on previous query execution times.

- Use the BigQuery command-line tool's `bq show` command to inspect the query.

### Question 26 Skipped

^

In the `europe-central2-a` zone, your team has an application server running on Compute Engine. What should you do to ensure high availability and replicate this server to `europe-central2-b` zone in as few steps as possible?

- You should use `gcloud` tool to copy the drive to the `europe-central2-a` zone, move this disk to `europe-central2-b` and finally create a new virtual machine with this disk.

- You should use `gcloud` tool to copy the drive to the `europe-central2-b` zone and then create a new virtual machine with this disk.

Correct answer



- You should create a snapshot from the disk and then create a disk from this snapshot in the `europe-central2-b` zone. Finally, create a new virtual machine with that disk.

- You should create a snapshot from the disk and then create a disk from this snapshot in the `europe-central2-a` zone, move this disk to `europe-central2-b` and finally create a new virtual machine with that disk.



### Question 27 Skipped



A mobile game developer wants to launch a new mobile game that will be available to users around the world. The game requires RDBMS for storing player profiles. As a cloud architect, which storage service should you recommend so they can scale to a global audience with minimal configuration updates?

Cloud SQL

Cloud Firestore

Correct answer



Cloud Spanner

Cloud Datastore

 **Question 28** Skipped ^

As a cloud architect, you are working with a client that needs an automated solution to back up their Compute Engine workloads. They want to ensure minimal downtime during the backup process and be able to restore their workloads to any point in the last 30 days. Which of the following solutions would be the most effective for this scenario?

- Use Cloud Datastore to keep the state of the Compute Engine instances.

Correct answer



- Use Persistent Disk Snapshots scheduled at regular intervals.

- Use Google Cloud's operations suite to monitor and backup Compute Engine instances.

- Use Cloud Storage to store periodic snapshots of the Compute Engine instances.

**Question 29 Skipped**

^

Your customer is planning to run a containerized workload in Google Cloud. The customer has the following requirements:

- the workload must be easily deployable and manageable
- the workload must be easily scalable and able to handle a large volume of data
- the workload must be able to handle failures and recover gracefully
- the workload must be easily integratable with other Google Cloud services

Which Google Cloud service should you recommend to meet these requirements?

Correct answer



Google Kubernetes Engine

Amazon EC2

Cloud Functions

Compute Engine



### Question 30 Skipped



Your company has a multi-tier web application running on Google Cloud. You've been tasked with planning for future improvements. User feedback indicates that the application occasionally experiences latency issues, which you suspect is due to increasing database load. What would be the best approach to address this issue and improve the application?

- Use Cloud Dataflow to preprocess and aggregate data before storing it in Cloud SQL.

Correct answer



- Migrate the database from Cloud SQL to Cloud Spanner to take advantage of horizontal scaling.

- Implement a caching layer using Cloud Memorystore to reduce database load.

- Increase the size of the Cloud SQL instance to accommodate the higher load.

● Question 31 Skipped ^

Your company wants to deploy a batch processing workload in Google Cloud using Apache Spark. The workload consists of a large number of compute-intensive tasks that can be executed in parallel. The workload must be scalable and cost-effective. How would you deploy the batch processing workload in Google Cloud to meet these requirements?

Correct answer



- Use Cloud Dataproc to run the batch processing tasks.
- Use a fleet of Compute Engine virtual machines to run the batch processing tasks.
- Use Cloud Batch to run the batch processing tasks.
- Use Cloud Functions to run the batch processing tasks.

**Question 32 Skipped**



Refer to the TerramEarth case study for this question: [https://services.google.com/fh/files/blogs/master\\_case\\_study\\_terramearth.pdf](https://services.google.com/fh/files/blogs/master_case_study_terramearth.pdf)

Upon evaluating TerramEarth's business requirements to minimize downtime, it was determined that a substantial amount of time-saving could be achieved by decreasing customers' wait time for parts. Consequently, the decision has been made to concentrate efforts on reducing the aggregate reporting time of three weeks. What changes to the company's processes would you suggest?

- Migrate from FTP to SFTP transport, develop machine learning analysis of metrics, and increase dealer local inventory by a fixed factor.

- Migrate from CSV to binary format, migrate from FTP to SFTP transport, and develop machine learning analysis of metrics.

- Migrate from FTP to streaming transport, migrate from CSV to binary format, and develop machine learning analysis of metrics.

Correct answer



- Increase fleet cellular connectivity to 80%, migrate from FTP to streaming transport, and develop machine learning analysis of metrics.

Question 33 Skipped ^

Your company offers a downloadable rendering software through its website, serving customers globally. To ensure optimal customer experience, you aim to minimize delays for all users while adhering to Google's recommended practices. What is the recommended approach for storing the files?

Correct answer



- Save the files in multiple Multi-Regional Cloud Storage buckets, one bucket per multi-region.

- Save the files in multiple Regional Cloud Storage buckets, one bucket per region.

- Save the files in multiple Regional Cloud Storage buckets, one bucket per zone per region.

- Save the files in a Multi-Regional Cloud Storage bucket.



### Question 34 Skipped



As a cloud architect, you are consulting for a manufacturing company implementing an IoT solution. They plan to collect real-time sensor data from their machinery across multiple plants, with high frequency. This data will be used to monitor machine health, and it's crucial to have low-latency access to the last hour of data. Which of the following would be the most effective way to use Google Cloud Bigtable in this scenario?

Correct answer



- Use Cloud Bigtable to store all sensor data, and use row keys that include the machine ID and timestamp to enable fast access to recent data.

- Use Cloud Bigtable to store all sensor data, and use row keys that include the sensor type and machine ID to enable fast access to recent data.

- Use Cloud Bigtable to store only the last hour of sensor data, and use a separate system for long-term storage.

- Use Cloud Bigtable to store all sensor data, but partition the data by machine ID and timestamp.

● Question 35 Skipped ^

As a cloud architect, you are working with a client who wants to migrate a large number of files from their on-premises data center to a Cloud Storage bucket. The client's on-premises data center is located in a region with a slow and unstable network connection. They would like the transfer to happen as quickly as possible, while also ensuring that the transfer can continue from where it left off in case of network interruptions. Which gsutil command option should you recommend for this scenario?

  -r

Correct answer

  -m  -c  -o

● **Question 36** Skipped ^

A SaaS solution for enterprise customers needs to be updated. Many components of the service are stateful and the system wasn't designed to allow incremental rollout of new code. The entire environment has to be running the same version of the deployed code. What implementation strategy would you recommend?

**Canary deployment strategy**

**A/B deployment strategy**

**Rolling deployment strategy**

Correct answer



**Blue/Green deployment strategy**

● Question 37 Skipped

^

A large retail company wants to launch a new personalized shopping experience for customers using an e-commerce platform. The platform should allow customers to create a profile, save their preferences and purchase history, and receive personalized product recommendations in real-time. The platform should also provide real-time analytics and reporting on customer behavior and purchasing patterns. Which of the following solutions would be the most effective for meeting these requirements?

Correct answer



**Use BigQuery for storing customer data, Cloud Dataflow for processing customer**

- preferences and purchase history, and Cloud Machine Learning Engine for generating product recommendations. Use Cloud Dataproc for analytics and reporting.**

**Use Cloud SQL for storing customer data, Cloud Functions for processing customer**

- preferences and purchase history, and Cloud Machine Learning Engine for generating product recommendations. Use Cloud Data Studio for analytics and reporting.**

**Use Cloud Datastore for storing customer data, Cloud Functions for processing**

- customer preferences and purchase history, and Cloud Machine Learning Engine for generating product recommendations. Use Cloud Data Studio for analytics and reporting.**

**Use Cloud Firestore for storing customer data, Cloud Functions for processing**

- customer preferences and purchase history, and Cloud AI Platform for generating product recommendations. Use Cloud Data Studio for analytics and reporting.**

● **Question 38 Skipped**

^

Your company offers online services that collect data about users and operates in North America. You want to start a business in Europe. What regulations must your company meet?

COPPA

SOX

HIPAA/HITECH

**Correct answer**



GDPR

● **Question 39 Skipped**

As a cloud architect, you are helping an organization set up a secure, low-latency network connection between their Compute Engine instances in the `us-central1` region and their local data center. They want to ensure the connection is private and the data isn't exposed over the public internet. Which of the following options should you recommend?

- Use a NAT gateway to enable communication between the Compute Engine instances and the local data center.

- Set up VPN Gateway on Google Cloud and enable IPsec for secure, encrypted communication.

Correct answer



- Set up Dedicated Interconnect at a colocation facility near their data center to connect directly to Google's network.

- Set up a VPC Network Peering between their Google Cloud VPC and their local data center.

#### Question 40 Skipped

^

A company wants to build a data processing pipeline that ingests data from various sources, transforms the data, and then stores the processed data in a data warehouse for analysis. The pipeline should be scalable, easy to manage, and able to handle changes in data sources and data processing requirements. How would you design this data processing pipeline using Google Cloud services?

Correct answer



- Use Cloud Dataflow to ingest data from the various sources, Cloud Dataproc to transform the data, and BigQuery to store the processed data.

- Use Cloud Functions to ingest data from the various sources, Cloud Dataflow to transform and store the data.

- Use Cloud Storage to ingest data from the various sources, Cloud Dataproc to transform the data, and BigQuery to store the processed data.

- Use Cloud Pub/Sub to ingest data from the various sources, Cloud Dataproc to transform and store the data.

**Question 41 Skipped**

^

Your company is developing a new IoT application that is expected to produce massive amounts of data for real-time analytics and future predictive models. Currently, the IoT devices publish data to a Pub/Sub topic, which then triggers a Cloud Function to store the data in BigQuery. What can be done to enhance this solution as data volume increases?

Use Firestore to store IoT data because of its real-time capabilities.

Migrate from BigQuery to Cloud SQL for data storage.

Correct answer



Migrate the data processing from Cloud Function to Dataflow to better manage streaming data.

Increase the memory and CPU allocated to the Cloud Function to process more data.

Question 42 Skipped ^

When migrating a legacy application to Google Cloud Platform (GCP), which of the following strategies is most likely to result in the quickest and least disruptive migration, while also taking into account future scalability needs?

Correct answer



- Lift and shift the application to Compute Engine Virtual Machines (VMs) and use Cloud Storage for data storage.
- Re-architect the entire application to use Google Cloud Platform services and design patterns.
- Use App Engine to host the application and store data in Cloud Datastore or Cloud Firestore.
- Use the Virtual Private Cloud (VPC) to create a dedicated network for the application and its components, and deploy the components to Compute Engine instances.

### Question 43 Skipped

^

You are a cloud architect tasked with architecting a data storage solution for a media company. The company has the following data storage requirements:

- store large media files that are regularly accessed for the first month
- archive media files that have not been accessed for over a year
- ensure redundancy and high availability
- keep costs optimized based on the frequency of data access

Which combination of storage classes should be used for Google Cloud Storage to meet these requirements?

Use Multi-Regional Storage for the large media files and Coldline Storage for archiving files not accessed in over a year.

Use Nearline Storage for all media files and enable Object Lifecycle Management to change the storage class to Coldline for files not accessed in over a year.

Correct answer



Use Standard Storage for the large media files and enable Object Lifecycle Management to change the storage class to Archive for files not accessed in over a year.

Use Standard Storage for all media files and enable Object Lifecycle Management to change the storage class to Nearline for files not accessed in over a year.

● **Question 44** Skipped ^

As a cloud architect, you're helping a large organization configure IAM roles for their security team. The team needs to audit corporate applications deployed on GCP, identify security risks, and recommend improvements, but they should not be allowed to modify resources. Which of the following is the most suitable IAM role to assign to this team?

Security Reviewer role

Project Editor role

Project Viewer role

Correct answer



Security Auditor role

Security Admin role

Question 45 Skipped ^

As a cloud architect, you need to plan your enterprise network in Google Cloud. Why should enterprises use custom VPC networks rather than the default network? (select 2)

Because in custom mode you automatically get some pre-populated firewall rules.

Because in custom mode you cannot choose unique, descriptive names for custom mode subnets.

Correct selection



Custom VPC networks integrate better with existing IP address management

schemes. The default networks use the same set of internal IP ranges. IP ranges might overlap when connected with your on-premises corporate networks.

Correct selection

Because you cannot connect two auto mode VPC networks to each other using VPC Network Peering because their subnets use identical primary IP ranges.

 **Question 46** Skipped 

You are configuring a fleet of Compute Engine instances that need to communicate with each other for data processing tasks. These instances need to access a Google Cloud Storage bucket. You want to optimize network throughput and costs. What should you consider?

**Correct answer**



- Place instances and Cloud Storage bucket in the same region.**
- Use Compute Engine instances with higher vCPU count and memory to increase network performance.
- Use dedicated Interconnect to connect instances and Cloud Storage.
- Enable Cloud CDN on Compute Engine instances.

● Question 47 Skipped

^

As a cloud architect, you are working with a large enterprise customer who is storing millions of daily logs in a Cloud Storage bucket. To manage costs and ensure data is retained appropriately, they want to automatically move files older than 30 days to Nearline storage and delete any files older than 365 days. Which of the following is the best approach to meet these requirements?

Correct answer



- Use the lifecycle management feature of Cloud Storage to move older files to Nearline and delete files older than 365 days.
- Use gsutil to move older files to Nearline storage and delete files older than 365 days manually.
- Use Cloud Storage Transfer service to move files to Nearline storage and lifecycle management to delete files older than 365 days.
- Use gsutil to periodically move all files to Nearline storage and use lifecycle management to delete files older than 365 days.

 **Question 48** Skipped 

You are designing a cloud-based architecture for a company that requires high availability and fault tolerance for its critical application. Which of the following options provides the best solution?

- Use a single instance virtual machine in one region with an attached disk.

Correct answer



- Use multiple instance virtual machines across different regions with a load balancer.

- Use a container-based approach with Kubernetes.

- Use a serverless computing approach with AWS Lambda.

● **Question 49** Skipped ^

As a cloud architect, you are working with an organization that wants to have granular control over its Google Cloud resources. The organization has multiple projects, each managed by different teams and belonging to different departments. The company wants to enforce the principle of least privilege and ensure that the resources are only accessible to the team that needs them, while also keeping the management simple and efficient. Which of the following approaches should you recommend?

Correct answer



- Create Google Groups for each team, assign roles to the groups at the project level, and add the team members to the appropriate groups.**
- Assign roles to the departments at the organization level and let the departments manage their individual teams' access.
- Assign individual roles to each team member at the project level.
- Assign roles to the teams at the organization level.

● Question 50 Skipped

^

Your company is building a multi-tier web application in Google Cloud. The application has a web frontend, an application backend, and a database backend. The database backend is critical to the operation of the application and must be highly available and scalable. How would you design the database backend to meet these requirements?

- Use a Cloud SQL read replica for the database backend and configure automatic failover.

- Use a single Bigtable instance for the database backend.

- Use a single Cloud SQL instance for the database backend.

Correct answer



- Use a Cloud Spanner instance for the database backend.

Question 51 Skipped ^

A financial services company wants to implement a secure and scalable solution for processing large volumes of financial transactions. The solution should also provide real-time visibility into transaction status, and enable employees to resolve issues and perform auditing tasks. The solution should be fully managed, and minimize the need for additional hardware and software investments. Which of the following options would be the most effective approach to meet these requirements?

Correct answer



- Use Cloud Spanner for storing transaction data and Cloud Dataflow for processing transactions in real-time. Use Cloud Pub/Sub for real-time updates on transaction status and Cloud IAM for controlling access to the data.**

- Use Cloud SQL for storing transaction data and Cloud Functions for processing transactions in real-time. Use Cloud Pub/Sub for real-time updates on transaction status and Cloud IAM for controlling access to the data.

- Use Cloud Datastore for storing transaction data and Cloud Functions for processing transactions in real-time. Use Cloud Pub/Sub for real-time updates on transaction status and Cloud IAM for controlling access to the data.

- Use Cloud Bigtable for storing transaction data and Cloud Functions for processing transactions in real-time. Use Cloud Pub/Sub for real-time updates on transaction status and Cloud IAM for controlling access to the data.

Question 52 Skipped ^

As a cloud architect, you are working with a global company that uses Cloud Storage for data storage. The company has a large number of contractors who need to upload data to a specific Cloud Storage bucket, but they should not have any other access rights to the data or the other resources. At the same time, the data engineers should have the ability to manage all Cloud Storage resources. What IAM roles should be assigned to the contractors and the data engineers?

- Assign roles/storage.objectAdmin to contractors and roles/storage.admin to data engineers.

- Assign roles/storage.objectCreator to contractors and roles/storage.objectAdmin to data engineers.

Correct answer



- Assign roles/storage.objectCreator to contractors and roles/storage.admin to data engineers.

- Assign roles/storage.objectViewer to contractors and roles/storage.objectAdmin to data engineers.

### Question 53 Skipped

As a cloud architect, you are working with a media company that runs a complex application on a Kubernetes cluster in Google Kubernetes Engine (GKE). After evaluating the application's resource usage over the past few months, you've determined that the cluster is over-provisioned and you need to resize it to save costs. The application is stateless, and there are no user sessions to maintain. The application is also globally distributed and must maintain high availability during the resizing process. What is the best approach to resize the cluster?

- Delete the existing cluster and create a new one with fewer nodes.

Correct answer



- Use the `gcloud container clusters resize` command to reduce the number of nodes in the cluster.

- Manually remove nodes from the existing cluster until you reach the desired size.

- Adjust the number of vCPUs and memory allocated to each node in the cluster.

**Question 54 Skipped**

^

Your customer is planning to run a data processing pipeline that ingests, processes, and stores large amounts of data in Google Cloud. The customer has the following requirements:

- the pipeline must be easily scalable and able to handle a large volume of data
- the pipeline must be able to handle failures and recover gracefully
- the pipeline must be easily integratable with other Google Cloud services
- the pipeline must be cost-effective

Which Google Cloud service should you recommend to meet these requirements?

Cloud BigQuery

Cloud Dataproc

Cloud Pub/Sub

**Correct answer**



Cloud Dataflow

● Question 55 Skipped ^

When deploying your application to App Engine, your development team aims to scale the number of instances in response to the request rate. It is necessary to maintain a minimum of five idle instances at all times. Which scaling method should they employ?

- Basic Scaling with `min_instances` set to 5.

Correct answer



- Automatic Scaling with `min_idle_instances` set to 5.

- Basic Scaling with `max_instances` set to 5.

- Manual Scaling with 5 instances.

● **Question 56 Skipped**

^

As a cloud architect, you're working with an organization that operates a complex microservices application on Google Kubernetes Engine (GKE). They want to utilize Service Mesh visualization in the Google Cloud Console to gain insights into their services' performance and interactions. However, after setting up their environment with Istio, they are unable to see any traffic flow between services. Which of the following could be a possible reason and the appropriate fix?

- The Service Mesh visualization doesn't support GKE. They should use Stackdriver for visualizing service interactions.

- Istio is not installed correctly. They should reinstall it and restart their services.

Correct answer



- The required Envoy proxy sidecar containers might not have been injected into each relevant Kubernetes pod. They should ensure automatic or manual sidecar injection is configured.

- Service Mesh visualization does not support microservices. They should refactor their application into a monolith.

● **Question 57 Skipped**

^

A global media company wants to implement a video streaming platform that can handle high traffic and deliver high-quality video to users. The platform should also be able to provide real-time analytics on video viewing trends and user engagement. The company wants to use the Google Cloud Platform for this implementation. Which of the following options would be the best solution for this requirement?

- Use Cloud Storage to store video data, Cloud Functions to process video data, and Cloud SQL to store analytics data.

Correct answer



- Use Cloud CDN to deliver high-quality video to users, Cloud Storage to store video data, and Cloud Bigtable to store video analytics data.

- Use Cloud Video Intelligence API to extract insights from video data, Cloud Pub/Sub to ingest video data, and Cloud Dataflow to process video analytics.

- Use Cloud Pub/Sub to ingest video data, BigQuery to store video data, and Cloud Dataflow to process video analytics.

Question 58 Skipped ^

As a cloud architect, you are assisting a client with an e-commerce application hosted on Google Cloud Platform (GCP). The client's application is expected to receive an increased volume of traffic during an upcoming sale. They want to distribute incoming traffic across multiple Compute Engine instances using Cloud Load Balancer. Additionally, they want to ensure that if any instance is unresponsive or unable to handle requests, it is automatically removed from the pool of available instances. Which of the following configurations would best achieve this?

- Configure a Network Load Balancer and set up a TCP health check.

Correct answer



- Configure an HTTP(S) Load Balancer with global backend services and set up an HTTP health check.

- Configure a Network Load Balancer and set up an HTTPS health check.

- Configure a TCP/SSL Proxy Load Balancer and set up a TCP health check.

**Question 59 Skipped**

^

As a cloud architect, you are working with an organization that frequently provisions and de-provisions a set of Google Cloud resources for temporary projects. They are looking for an automated and reproducible way to manage this. Which of the following approaches best utilizes Google Cloud Deployment Manager for this purpose?

- Use the Deployment Manager API to manually create and delete resources as needed.
- Use Deployment Manager to automate the process of manually creating and deleting resources in the Google Cloud Console.
- Use Deployment Manager to periodically check for unused resources and delete them.
- Create a Deployment Manager configuration file for the required resources, and use Deployment Manager to create and delete these resources as needed.

Correct answer



**Question 60 Skipped**

^

A large financial services company wants to build a new trading platform that can process millions of transactions per second. The platform should be able to handle high levels of volatility, and provide low latency data access for real-time decision making. Which of the following options would be the best solution for this requirement?

Correct answer



- Use Cloud Pub/Sub for data streaming and BigTable for data storage.
- Use Cloud Spanner for data storage and Cloud Dataflow for data processing.
- Use BigQuery for data storage and processing.
- Use Cloud Datastore for data storage and Cloud Functions for data processing.



### Question 61 Skipped



Refer to the Helicopter Racing League (HRL) case study for this question: [https://services.google.com/fh/files/blogs/master\\_case\\_study\\_helicopter\\_racing\\_league.pdf](https://services.google.com/fh/files/blogs/master_case_study_helicopter_racing_league.pdf)

The Helicopter Racing League (HRL) is seeking your assistance in expanding the reach of their existing recorded video content to attract new fans in emerging regions. In light of the business and technical requirements of HRL, what actions do you need to take?

- You should replicate the video content in Google Kubernetes Engine clusters in regions close to the fans.

- You should use Apigee Edge to cache the video content from HRL's existing public cloud provider.

- You should serve the video content directly from a multi-region Cloud Storage bucket.

Correct answer



- You should use Cloud CDN to cache the video content from HRL's existing public cloud provider.

● **Question 62 Skipped**

^

As a cloud architect, you need to establish connection between your on-premises network and Google Cloud. Your company will need 5 Gbps of bandwidth in total between the on-premises data center and Google Cloud. The traffic may be split between multiple connections. How many VPN endpoints will you need?

Correct answer



2

4

1

3

● Question 63 Skipped ^

As a cloud architect, you're helping a large organization deploy an update to their mission-critical application on App Engine. The application is used 24/7 by users worldwide, and it is crucial that the update does not affect application availability or user experience. Which of the following deployment strategies should you recommend?

- Blue-Green Deployment: Create a new version of the application, migrate all traffic to the new version, and delete the old version.

- Rolling Update: Update the application code in the existing version and restart the application.

Correct answer



- Canary Deployment: Create a new version of the application and gradually migrate user traffic from the old version to the new version.

- Big Bang Deployment: Replace the existing application with the new version in a single operation.

#### Question 64 Skipped



A large healthcare organization is looking to build a secure and scalable platform for storing and analyzing medical records. The platform must meet the following requirements:

- support high volumes of medical records with low latency
- ensure secure storage and processing of sensitive medical information
- enable real-time data analysis and reporting on patient health and treatment outcomes
- minimize downtime during maintenance and upgrades
- minimize costs while still providing high performance

Which solution would you recommend to meet these requirements?

- Implementing a custom-built solution using Cloud Pub/Sub for real-time data processing, Bigtable for data storage, and Google Kubernetes Engine (GKE) for deployment and scaling.

- Correct answer 
- Implementing a managed solution using Cloud Healthcare API for data analysis, Cloud Storage for data storage, and Cloud Load Balancing for high-availability data access.

- Implementing a serverless solution using Cloud Functions for data processing, Cloud Firestore for data storage, and Cloud Pub/Sub for real-time data processing.

- Implementing a hybrid solution using Compute Engine for data processing, Cloud SQL for data storage, and BigQuery for real-time analytics.

### Question 65 Skipped

^

A large media company is looking to build a secure and scalable platform for hosting and streaming video content. The platform must meet the following requirements:

- support high volumes of concurrent video streams with low latency
- ensure secure storage and processing of sensitive video content
- enable real-time analytics and reporting on video usage and performance
- minimize downtime during maintenance and upgrades
- minimize costs while still providing high performance

Which solution would you recommend to meet these requirements?

Correct answer



- Implementing a managed solution using Video Intelligence API for video analysis, Cloud Storage for video storage, and Cloud Load Balancing for high-availability video streaming.**

- Implementing a custom-built solution using Cloud Pub/Sub for real-time data processing, Bigtable for data storage, and Google Kubernetes Engine (GKE) for deployment and scaling.**

- Implementing a hybrid solution using Compute Engine for video processing, Cloud Storage for video storage, and BigQuery for real-time analytics.**

- Implementing a serverless solution using Cloud Functions for data processing, Cloud Firestore for data storage, and Cloud Pub/Sub for real-time data processing.**

● **Question 66** Skipped ^

In your role as a cloud architect, you're preparing to transition your on-site data warehouse to Google Cloud, utilizing BigQuery. You have to create a presentation for the management team to explain the cost structure of BigQuery. There are two primary elements to consider when it comes to the pricing of BigQuery. Select the costs you incur when using BigQuery.

- The cost of Identity and Access Management (IAM pricing).

Correct selection



- The cost to store data that you load into BigQuery (storage pricing).

- The cost of viewing table schemas (schema pricing).

Correct selection

- The cost to process queries (analysis pricing).

● Question 67 Skipped ^

You are designing a Google Cloud architecture for an organization that operates in a heavily regulated industry. It needs to meet strict compliance requirements regarding data storage and access. Which approach should you use?

Use Cloud Audit Logs to monitor and audit data access.

Enable VPC Service Controls to limit data exfiltration.

Correct answer



Use customer-managed encryption keys (CMEK) and set up access controls for the keys.

Encrypt all data at rest using Google-managed encryption keys.

**Question 68 Skipped**

^

A marketing company creates a lot of landing pages (static websites). As a cloud architect, which storage service would you recommend in this case to minimize costs?

Cloud Datastore

Cloud Endpoints

Compute Engine + Persistent Disk

Cloud SDK

**Correct answer**



Cloud Storage

Question 69 Skipped ^

Refer to the Mountkirk Games case study for this question: [https://services.google.com/fh/files/blogs/master\\_case\\_study\\_mountkirk\\_games.pdf](https://services.google.com/fh/files/blogs/master_case_study_mountkirk_games.pdf)

As the Data Compliance Officer for Mountkirk Games, your responsibility is to safeguard customers' personally identifiable information (PII). The company seeks to generate anonymized usage reports for its new game and implement a data deletion policy for PII after a designated timeframe. Your role is to ensure compliance while considering the business and technical requirements with minimal cost implications. What course of action would you recommend?

Correct answer



- You should archive audit logs in BigQuery, and generate reports using Google Data Studio.

- You should write a Cloud Logging filter to export specific date ranges to Pub/Sub.

- You should archive audit logs in Cloud Storage, and manually generate reports.

- You should archive user logs on a locally attached persistent disk, and cat them to a text file for auditing.

**Question 70 Skipped**

^

Your company has a massive amount of data (approx 50 TB) in its on-premises data center, which needs to be transferred to Cloud Storage for further analysis and machine learning processing. Due to security concerns, your company has decided not to use physical data transfer services like Transfer Appliance. Your company has a very high-speed internet connection, but the rate of data generation is also high, leading to a shrinking time window for data transfer. What would be the best way to migrate this data over the internet in a reasonable amount of time?

- Use a combination of Google Cloud's Data Transfer Service and gsutil for multi-threaded transfers.
- Use gsutil with multi-threaded/multi-processing options to copy data directly into Cloud Storage.

Correct answer



- Use Storage Transfer Service with Cloud VPN for secure data transfer.
- Use Cloud Dataflow to transfer the data.